



RAPS v0.7.0 Repository Audit

1. Missing or Incomplete Features (vs. Project Goals)

ACC Modules: While RAPS covers Autodesk Construction Cloud (ACC) Issues, several ACC services mentioned in the design are not yet implemented. The code and roadmap reference support for RFIs, Assets, Submittals, and Checklists, but these modules are only placeholders as of v0.7.0 [1](#) [2](#). In other words, beyond **Issues**, the ACC domain is incomplete – features like RFI management or asset tracking are not available yet.

Plugin/Extension System: A plugin/extension architecture is sketched out but not finished. The code includes a **Plugin** system (allowing external `raps-<name>` executables, hooks, and aliases) with a config file format defined [3](#) [4](#). However, this is a **work-in-progress**: there are no user commands or docs to actually manage or use plugins in v0.7.0, and the changelog explicitly lists the plugin system as an “architecture (placeholder)” [5](#). Thus, extensibility via plugins is not functional yet.

Other APS APIs: The CLI aims to be “comprehensive” for Autodesk Platform Services, and it does cover the major APIs (Authentication, OSS, Model Derivative, Data Management, Webhooks, Design Automation, Issues, Reality Capture). No obvious entire API is missing from the core APS suite. However, within those domains some **advanced capabilities** might be lacking. For example, the **pipeline automation** (new in 0.7.0) is a powerful addition, but it’s sequential-only – there’s no support for parallel step execution or looping (beyond what a user scripts manually). Similarly, **Data Management** commands cover browsing and basic item operations, but not more complex project admin tasks. These are not necessarily promised in the design, but could be future feature gaps as usage grows.

Machine-Readable Output Formats: Early roadmaps show an intent for consistent machine outputs (JSON/YAML) for all commands [6](#). By v0.7.0, JSON output is supported widely and **YAML output was added** [7](#). One minor omission is that not all commands’ outputs are documented in a structured schema (making automation fragile if formats change). The CLI does not currently provide a formal schema or versioning for JSON outputs, which could be considered a missing piece if strict CI integration is a goal.

In summary, **most core features are implemented**, but **extensions and a few APS niches remain incomplete**. Support for additional ACC features (RFIs, assets, etc.) is clearly intended but not delivered [2](#), and the plugin system is in prototype stage. These gaps should be addressed to fulfill the project’s goal of a “comprehensive” APS CLI.

2. Documentation Gaps or Deficiencies

Breadth of Documentation: The project’s documentation is generally **extensive and well-organized**. There is a detailed README and a full docs site (via MkDocs) covering installation, configuration, and command usage. For example, the docs include step-by-step “Quick Start” and workflow examples [8](#) [9](#), which is excellent for new users. The README itself highlights all major features and provides installation

instructions, and there are dedicated pages for each command group (Auth, Buckets, Objects, Translate, etc.), complete with examples and option descriptions.

Despite this strength, a few documentation gaps exist:

- **Plugin System Documentation:** As noted, the plugin/extension feature isn't user-ready, and correspondingly there is no user documentation for it. The only hints are in code comments ¹⁰. If this feature is part of the design vision, its absence in the docs could confuse users who see mentions in changelogs or code. Once the plugin system is ready, comprehensive docs (how to create plugins, config file format, security considerations) need to be added.
- **Advanced Configuration & Tuning:** While the **configuration** page covers profiles and env vars well, there is minimal mention of tuning performance parameters (e.g. the `--concurrency` flag default and limits). For example, the CLI supports `--timeout` and `--concurrency` as of v0.5.0 ¹¹, but the documentation could better explain how to use these for large batch operations or how to adjust them for different network conditions. Similarly, **pagination behavior** is not documented – if a project has hundreds of items, does `raps item list` retrieve all? Currently, it's not clear in docs; users might hit page-size limits of the API without realizing. Documenting how the CLI handles pagination or providing options to fetch more results would improve clarity.
- **Internal/Developer Docs:** For end-users, docs are good, but for contributors there is limited high-level architecture documentation. There is a "Feature Overview" with a Mermaid diagram mapping commands to APIs ¹² ¹³, which helps understand the scope. However, there's no narrative design document explaining module structure or how error handling and logging are implemented. Code comments exist at module level (e.g. in `src/api/mod.rs` and for each client), but many functions lack rustdoc comments. Enhancing code comments (especially for complex sections like the pipeline executor or retry logic) would aid maintainability. The contributing guide covers the process but not the architecture. This is a minor gap, but important for an open-source "production-grade" project.
- **Known Limitations:** The documentation does not call out current limitations or unsupported scenarios. For example, it might be worth explicitly noting that certain APS services (like ACC checklists or assets) are not yet supported to set user expectations. Likewise, advising that the CLI is in beta (0.x) and interfaces may change would be prudent in the README. Currently, the README markets the tool as comprehensive, which is mostly true, but a "roadmap" or "limitations" section in the docs would add transparency. (There is a `ROADMAP.md` in the repo for developers ¹⁴, but an end-user might not see that.)

In summary, the documentation is strong on usage (installation, commands, examples) ⁸ ¹⁵, but could be improved by documenting **what's not there** (features or edge cases not yet handled) and by adding more **developer-facing context** (architecture and inline code docs). Keeping the docs in sync with new features (which the project has done well so far ¹⁶) remains an ongoing need.

3. Roadmap Items Unimplemented as of v0.7.0

The repository's roadmap (covering v0.4 through v0.6 milestones) has been largely fulfilled by v0.7.0 – in fact, an **implementation status** report shows all planned high-priority features up to 0.6 marked as

completed [17](#) [18](#). That includes CI/CD friendly output and flags, profile management, headless auth, retry logic, and supply-chain enhancements. As a result, few items from the official roadmap remain truly “unimplemented” by 0.7.0. Notable exceptions and post-0.6 plans include:

- **Plugin/Extension System:** As discussed, the roadmap didn’t explicitly list this in v0.4–0.6, but the v0.7.0 changelog identifies the plugin system as a new architecture goal [5](#). This is currently incomplete (placeholder in code). Implementing this (with proper security and documentation) is a future item.
- **ACC Extended Modules:** The idea of expanding to more ACC modules (beyond issues) is mentioned in v0.7.0 notes [19](#). This presumably covers Assets, Checklists, Submittals, etc. These remain unimplemented at 0.7.0, so they are effectively roadmap items for a future version (perhaps v0.8 or v1.0).
- **Enhanced Error Context:** The early roadmap review suggested improving error messages with more context [20](#). By 0.7.0 some progress was made (the changelog notes “*human-readable explanations*” for errors [21](#)), but this is an area that can always be refined. It might not be fully “done” if the aim is to have extremely clear suggestions for every common error code.
- **Testing & CI Improvements:** Not explicitly a roadmap item, but often part of “production readiness” – things like increasing test coverage or adding integration tests were not listed in the roadmap. If there’s a roadmap beyond 0.7.0, it should likely include **QA-related tasks** (more on this in section 5).

Apart from these, most of the v0.4–0.6 roadmap issues are implemented as features in 0.7.0. To double-check: global `--output` with JSON/YAML is done [7](#), standardized exit codes are done [22](#), logging flags are done [23](#), non-interactive mode is implemented for all key commands [24](#) [25](#), profile management and keychain integration are done [26](#), device-code and token auth done [27](#), retry/backoff and timeouts done [28](#), proxy support documented [29](#), checksums and SBOM done [30](#), changelog and issue templates done [31](#). The **roadmap items that remain open are essentially those introduced in or after v0.7.0** – primarily the extension/plugin framework and the remaining APS/ACC coverage. Going forward, formalizing a roadmap to 1.0 (covering these items) would be wise.

4. Open Issues / PRs Indicating Bugs or Enhancements

As of the v0.7.0 release, there do not appear to be any critical open bug reports in the repository. The development issues tracked in this project were mostly roadmap features and enhancements, which have been closed out as they were implemented (15 issues covered v0.4–0.6 milestones [32](#), all resolved in subsequent releases). A search through the issue tracker shows feature tasks and chores (e.g. adding templates, timeouts, SBOM generation) which are closed, corresponding to completed work in the changelog [33](#) [34](#). No severe bug reports (crashes, data loss, etc.) are documented, suggesting a stable tool for the implemented features.

However, a few **enhancement requests / known limitations** can be inferred:

- There was a PR or issue about “*pipeable output formats*” (making output easily pipeable to other programs) [35](#). This likely relates to ensuring quiet mode, no-color output, etc., which were

implemented in v0.4.0. The existence of that PR suggests a focus on polishing output for scripting; it's presumably merged, but any related open points (like ensuring absolutely no extraneous output in `--quiet` mode) should be verified.

- **No interactive issues open:** Since non-interactive mode has been implemented, there are no open issues about automation blocking prompts. But the team should remain vigilant for any command that still might require input in edge cases (this might pop up via user issues if someone finds a missed prompt).

- **Potential Known Limitations:** Even if not filed as issues, a few limitations are known from design:

- *Pagination/large data handling:* As noted, the CLI might not yet handle multi-page results from the APS APIs. This could be considered a known limitation (the user may need to manually supply `--limit` flags or the CLI only fetches the first page). This isn't recorded as an issue, but it's a known aspect to improve for scalability.
- *Plugin system:* The incomplete plugin/alias feature is a known limitation (not filed as a bug, since it's more of a future enhancement).
- *ACC domain coverage:* It's known that support for other ACC modules isn't there yet – essentially a feature gap, as discussed.
- *Stability under stress:* No issues have been filed about memory leaks or performance, but real-world usage might uncover some (e.g. uploading extremely large files or thousands of items in one go). There's no evidence of such bugs in the tracker, but that's an area to monitor as adoption grows.

In summary, **open issues are minimal** – the project appears well-triaged and most known tasks were resolved by 0.7.0. The "issues" that do exist are mostly planned enhancements (like the remaining features) rather than bug fixes. This is a good sign, but also means the team should seek more feedback from users to catch any subtle bugs not yet reported.

5. Assessment for Production-Grade Readiness and Recommendations

RAPS is **quite advanced for a 0.7.0 CLI** – it implements a broad set of features with attention to CI/CD usage, error codes, and security. To elevate it to a truly production-grade tool (approaching a 1.0 release suitable for enterprise use), the following areas should be addressed:

Stability & Reliability

RAPS has already introduced standardized exit codes and a global non-interactive mode to make scripts reliable ³⁶. It also includes retry and timeout logic to handle API flakiness ²⁸. These are critical for stability. Moving forward, the project should:

- **Hardening through Testing:** Perform extensive testing of all commands in various scenarios (network outages, API rate limiting, invalid inputs) to ensure the CLI behaves predictably (retries when it should, fails fast when it should). For example, simulate a 429 or 500 error from APS and verify the retry/backoff works and eventually returns exit code 5 (remote error) as expected ³⁷. Any

remaining inconsistent behaviors (if any command doesn't conform to the standardized exit codes or output formats) should be fixed.

- **Graceful Degradation:** Ensure that if a particular APS service is down or returns an error, the CLI reports it clearly and doesn't crash. The error handling improvements in 0.7.0 (human-readable messages and suggestions) are great ²¹; continuing to expand those (maybe mapping common error codes to advice, e.g. 401 -> "authentication expired, please login again") will improve reliability in production use.
- **State Consistency:** Consider transaction-like behavior for multi-step operations. For instance, the pipeline feature executes sequences – if one step fails and `continue_on_error` is false, the pipeline stops. That's fine, but ensure partial results are still usable or recoverable. Similarly, for batch operations, if one file fails to upload, does the CLI continue or abort? The design should be clear and documented for these scenarios, so users can rely on consistent behavior.
- **Robustness in CI:** Since one goal is CI/CD usage, verify that the CLI can run unattended for long jobs. Memory leaks or resource exhaustion would be unacceptable in production – running a long translation or processing hundreds of files should not degrade the system. This may require profiling under stress to ensure reliability in long-running processes.

Overall, RAPS is **already quite stable**, but the above steps will solidify confidence that it won't hang, crash, or mis-report status even under heavy use or adverse conditions.

Performance & Scalability

Performance optimizations are already evident: for example, RAPS supports **multipart uploads and parallel transfers** to handle large files efficiently ³⁸. To be production-grade, consider:

- **Profiling and Optimization:** Measure performance for key operations (upload, download, listing large projects, running a full pipeline). Ensure that operations scale linearly and there are no obvious bottlenecks in code (e.g. inefficient JSON processing or unnecessary memory copies). The Rust implementation and use of `async/tokio` should help here, but real-world profiling might reveal tweaks (like buffer sizes or concurrency defaults) to improve throughput.
- **Scalability of Listings:** Introduce or document **pagination support**. In enterprise use, there may be hubs with hundreds of projects or projects with thousands of items. If the current CLI only fetches a default page (say 100 items), that's a limitation. Ideally, RAPS should either automatically paginate through all results (which might be implemented internally via looping API calls), or provide flags like `--all` or `--page-token` so advanced users can fetch everything. At minimum, document the limit and allow the `--limit` parameter if the API supports it. Not addressing this could lead to silent truncation of results in production, which is dangerous.
- **Memory Footprint:** Ensure that streaming is used for large file transfers (it likely is, using `reqwest` with streaming). Loading entire multi-GB files into memory would be a problem. The implementation should use backpressure or chunking (the presence of chunked upload suggests it does). Verifying that **download** uses streaming as well (not buffering entire files) would be prudent. If not already, implement streamed downloads for very large derivatives or attachments.

- **Concurrency Controls:** The CLI now has a `--concurrency` flag for batch ops (default 5) ³⁹. That's good, but in production one might want even finer control (e.g. separate thread pools per service, or preventing too many concurrent API calls that could hit rate limits). The current default of 5 parallel uploads is reasonable; just ensure this is easily configurable via a global config or environment (it appears to be CLI flag only – consider allowing a config default). Also, consider if certain operations should **auto-adjust concurrency** (for example, if a user tries `--parallel 50`, perhaps detect and warn if the system or API might not handle it).

In summary, the performance is likely adequate for moderate use, but for **enterprise-scale scenarios** the team should validate and possibly enhance pagination and resource handling. The goal is to ensure the CLI scales from a quick interactive tool to handling large batch jobs without hiccups.

Security Considerations

Security has been thoughtfully considered so far. Key points: OAuth token handling is central (with support for device code and external tokens for CI), and there's an opt-in OS keychain integration to securely store tokens ⁴⁰. Also, the release process includes publishing SHA256 checksums and even SBOM generation for supply-chain transparency ³⁴. To further strengthen security for production:

- **Credential Storage:** Consider making the OS keychain storage **enabled by default** on platforms where it's available (with a flag to opt-out, rather than opt-in). Currently, the default stores tokens in a local JSON file (likely under the user profile) – for production use, that file could be a target (e.g. multi-user machines or shared CI runners). Using the OS secure storage by default (when feasible) provides better security out-of-the-box. At minimum, document the security model clearly: e.g. *"Tokens are stored unencrypted in `~/.config/raps/tokens.json` unless you enable keychain storage via `RAPS_USE_KEYCHAIN=1`".* This transparency lets users make informed decisions.
- **Secret Handling in Logs:** It's great that debug logs redact secrets ⁴¹. The team should continue to audit logging statements to ensure **no sensitive data is ever printed in any mode** (except perhaps an explicit `inspect` command). This includes access tokens, refresh tokens, client secrets, etc. The `auth inspect-token` command shows token details but presumably not the full token (and warns on expiry) – this seems fine. Just ensure any error messages from APIs that might echo back sensitive info are also filtered (unlikely, but worth noting).
- **Dependency and Binary Security:** The SBOM is generated (CycloneDX per notes) and signatures can be applied to releases. For a production tool, ensure that **binaries are signed/notarized** for each platform. There is a script for signing Windows, Mac, Linux binaries ⁴² ⁴³ – using this in the release process means enterprise users can verify the authenticity of the CLI. Continue to publish checksums (which is already done ⁴⁴) and consider publishing GPG signatures of releases for users to verify. Also, keep dependencies up to date with security patches (Rust makes this relatively easy via `cargo audit` or dependabot). Regularly run a vulnerability scan (e.g. `cargo audit` in CI) to catch any known issues in dependencies (reqwest, etc.) and update promptly.
- **Plugin Security:** If the plugin system is to be enabled, security must be a top concern. Executing arbitrary external programs via a CLI plugin mechanism can be risky. The team should implement safeguards: e.g. only load plugins from trusted locations, perhaps require users to explicitly enable a discovered plugin (to avoid accidentally running a malicious `raps-foo` found in the PATH). They

might also want to sign plugins or at least warn about unsigned third-party code. Since this is future, it should be planned now: a production-grade CLI should not introduce vulnerabilities via extension points.

- **API Permissions:** Make sure the CLI only requests the scopes it needs. It already allows fine-grained scope selection on login (the interactive prompt for scopes) and has default scopes ⁴⁵ ⁴⁶. That's good security practice (principle of least privilege). In production environments, documentation should encourage using a separate APS application with only needed scopes for automation, to limit damage if a token is compromised. Possibly provide a way to configure default scopes in a profile, so organizations can enforce that only certain scopes are used.

Overall, the project shows **strong security awareness** (OAuth flows, keychain, checksums, SBOM) ³⁴. The recommendations are about making secure defaults and continuing vigilance so that as the project grows (especially with plugins), security doesn't backslide.

Error Handling & Logging

By v0.7.0, error handling is much improved: exit codes are standardized and documented, and the CLI has verbose and debug modes for insight ⁴⁷ ⁴⁸. Production-grade recommendations:

- **Consistent User Messages:** Ensure every user-facing error message is actionable. For example, if a command fails with "404 Not Found", the CLI might augment it with "*(The resource was not found – check the ID or your permissions)*". The changelog mentions "*contextual suggestions for common API errors*" were added ²¹ – continuing to enrich these will greatly help end-users and reduce support load. Particularly, handle auth errors gracefully (e.g. if token expired, suggest `raps auth login` again).
- **Logging to File:** In CI or server use, it might be useful to have an option to tee verbose logs to a file (for debugging after the fact). Currently, `--verbose` / `--debug` print to stderr/stdout. Consider adding an environment variable or config setting for a log file path, so that a detailed log can be kept without cluttering the console output. This can be as simple as writing debug logs to `~/config/raps/raps.log` (rotated by size perhaps). It's not strictly necessary, but many production tools support capturing logs for audit.
- **Structured Logging:** Another enhancement could be to allow logs in JSON (especially for debug mode). This ties into pipeability – if users want to aggregate logs from many runs, JSON output (for logs, not command results) could be helpful. For instance, a `RAPS_LOG_FORMAT=json` env var could make debug logs emit machine-parseable entries (with timestamp, level, message, maybe command context). This would complement the existing human-readable logs and is useful in complex CI systems.
- **Suppressing Output:** Verify that `--quiet` truly silences everything except essential output. In scripting, one often wants to capture just a result. Ensure that under `--quiet`, even things like progress bars or prompts are fully suppressed. If any third-party library prints unexpectedly (e.g. indicatif progress bar might need explicit disable), handle that in code. Production users will appreciate that quiet means quiet, and all messaging obeys the flags.

- **Error Propagation:** Make sure errors don't get swallowed. For example, if a pipeline step fails, the pipeline command should end with exit code 1 (indicating failure) unless `--continue-on-error` was used ⁴⁹. From the docs this appears to be the case. As an improvement, consider if multiple errors occur in a pipeline (with continue-on-error), perhaps at the end print a summary of which steps failed. Tiny usability enhancements like that can make a big difference in troubleshooting complex runs.

In summary, **logging and error handling are on the right track**. The focus should be on polish: making messages as clear as possible, and giving ops teams the tools needed (log files, structured logs) to diagnose issues in a production context.

Test Coverage & CI/CD

For a production-grade release, test coverage and continuous integration should be very robust:

- **Unit and Integration Tests:** The project currently has a suite of unit tests (73 tests by v0.7.0) ⁵⁰. It's important to ensure critical functionality is covered by tests. Increase the depth of testing for things like: token refresh logic, retry/backoff (simulate a sequence of failures and recovery), pipeline parsing and execution (test with sample YAML files), and commands that involve file I/O (maybe using temporary files to test upload/download in a local scenario if possible). Where real APS calls can't be made in CI (due to needing credentials), consider using APS sandbox or mock the HTTP calls (the design could inject a mock client). At the very least, integration tests for flows could be run against a development APS application in a controlled environment. This will catch breaking changes in APS APIs early as well.
- **Cross-Platform CI:** Ensure that CI runs on all target OSes (Windows, Linux, macOS) to catch platform-specific issues (path handling, TLS differences, etc.). The GitHub Actions badge indicates a CI workflow ⁵¹; that workflow should build and test on all major platforms. This is especially crucial for a CLI to be released as binaries.
- **Continuous Delivery:** The release process is automated by tagging (per CONTRIBUTING.md) ⁵², which is good. For production readiness, also set up CI jobs to automatically publish the crate to crates.io and attach binaries to GitHub Releases upon tagging. It appears this is already done (since installation via `cargo install raps` is documented ⁵³ and a 0.7.0 tag exists). Verify that the release artifacts include everything (binaries, checksums ⁴⁴, SBOM, documentation link).
- **CI Quality Gates:** Use CI to enforce code quality – e.g. continue running Clippy and formatting checks (as listed in CONTRIBUTING) on each PR. Possibly introduce a **coverage report** requirement (e.g. ensure coverage doesn't decrease). Also consider adding `cargo audit` to CI to flag dependency vulnerabilities. These gates ensure ongoing code health as contributions increase.
- **Real-World Testing:** Before declaring 1.0, it would be wise to conduct a beta program or dogfooding within Autodesk (if applicable) – run RAPS in real production scenarios (like nightly jobs or large migrations) to see if any issues surface. This kind of battle-testing often reveals corner cases that unit tests might miss (especially performance-related or memory issues).

By beefing up testing and CI, the maintainers can confidently assert the CLI is production-ready. It also helps prevent regressions: as new features like plugins or new ACC commands are added, the tests should grow to cover them, so nothing breaks existing functionality. Overall, a **strong CI/CD pipeline is in place**, but it can be expanded to include more automated quality checks that are typical for production-grade software.

Versioning and Release Process

For a production-grade tool, proper versioning and a disciplined release process are key:

- **Semantic Versioning:** The project already adheres to semver (0.x for now) and maintains a detailed **CHANGELOG** ⁵⁴. To reach production status, planning a **v1.0.0** release is important. This means deciding which features are “must-have” for 1.0 (likely everything on the roadmap plus stabilization), and then bumping to 1.0 with a guarantee of backward compatibility. Until now, as a 0.x project, breaking changes have been allowed (and indeed happened as features were added). With 1.0, communicate clearly that breaking changes will require a major version bump going forward (and honor that). The existence of the changelog and version timeline will help users follow along ⁵⁵.
- **Release Frequency and Support:** Determine a sensible release cadence. Production users appreciate predictability – for example, after 1.0, maybe use a monthly or quarterly release schedule for new features, with patch releases for bug fixes as needed. Also, consider whether you’ll support older versions (e.g. if a critical bug is found in 1.0 after 1.1 is out, will 1.0 get a patch or must users upgrade?). Given this is a CLI, likely encouraging upgrade to latest is fine, but in enterprise settings some may pin a version. Thus, maintaining a **stable branch or LTS** might be worth considering once the user base grows.
- **Automating Releases:** The current process (tag -> CI build -> publish) should be tested to ensure it works flawlessly. The inclusion of checksums and SBOM in 0.6.0 was great for release integrity ³⁴. For 1.0, also consider providing a **homebrew formula** or other package manager integration to ease installation. Production users on Linux might want a Debian/RPM package. These are nice-to-haves that can be added to the release process (perhaps generating Linux packages in CI, etc.). At minimum, the pre-built binaries and `cargo install` cover a lot of ground.
- **Documentation Versioning:** As the releases progress, ensure documentation is versioned or at least updated in sync. The README and docs currently track “latest” (with a badge linking to latest docs). For production, consider tagging documentation for each release (so users on an older version can reference the docs for that version if needed). This might involve maintaining a `docs-v1.0` branch or using GitHub Pages versioning.
- **Backward Compatibility Strategy:** Before 1.0, do a sweep of the CLI arguments and outputs to decide if any changes are needed. For instance, ensure naming is consistent (command and flag names should be logical and final). If any breaking tweaks are needed, do them now in 0.x. After 1.0, adopt a deprecation policy: e.g. if you need to rename a command or option, support the old one (with a warning) for a couple of releases. Communicate changes clearly via the changelog and perhaps console warnings. This way, users have confidence that scripts built around RAPS won’t suddenly break on upgrade.

The project already shows maturity in versioning (detailed changelog, respecting semver principles ⁵⁶). By formalizing the release process and backward compatibility commitments at 1.0, it will meet enterprise expectations.

User Configuration & Customization

RAPS provides flexible configuration: environment variables, `.env` files, and profile-based config are all supported ⁵⁷ ⁵⁸. To further improve:

- **Configuration Consistency:** Verify that *every configurable parameter* (timeouts, concurrency, output format, etc.) can be set via either a flag or a config file or env var. Currently, core settings like `client_id`, `secrets`, etc., can be set in profiles or env. The precedence rules (CLI > env > profile > default) are documented ⁵⁸ and should be kept consistent. One recommendation is to allow setting some of the global flags via env/profile as well (for example, a user might want `RAPS_OUTPUT=json` in their profile so that by default all commands output JSON unless overridden by a flag). If such features exist, document them; if not, consider adding for convenience.
- **Profiles and Contexts:** The profile system is a strong point, enabling multiple environment configs. For production use, ensure that profile commands are rock-solid (no chance of corrupting the profile store). Possibly add the ability to **merge or inherit profiles** (for example, a base profile for common settings and a child profile for specific overrides). This might help in complex use cases but could also overcomplicate – it's an idea to evaluate based on user feedback.
- **Customization Hooks:** Once the plugin/hook system is implemented, users will have more customization power (like running pre/post commands or adding aliases as in the plugin config example ⁵⁹ ⁴). This will be a boon for advanced users. The recommendation is to implement and test this thoroughly before 1.0, as it will allow power users to adapt RAPS to their workflows (making it truly “automatable” for pipelines). For example, an enterprise might want a pre-hook to refresh tokens or a post-hook to send a notification on job completion – such things can be done via the planned hook config. Ensuring this works and is secure will elevate the CLI's usefulness.
- **User Feedback Mechanisms:** Provide ways for users to know what's configurable. The docs do a good job listing env vars (like `APS_CLIENT_ID`, etc.) ⁶⁰ ⁶¹. Continue to update this as new config options arise (e.g., if a `RAPS_TIMEOUT` or `RAPS_CONCURRENCY` env var is supported, list it). Perhaps a `raps config show` command could dump all current config values and sources – this could help troubleshooting in production setups by showing exactly what config is in effect (given the layered precedence). This isn't present yet, but would be a nice diagnostic feature.

In short, the configuration system is already **flexible and well-documented**, supporting production scenarios like multiple environments and CI secrets. A few more enhancements (global defaults, hooks, config diagnostics) would make it even stronger.

Backward Compatibility Guarantees

Currently, as a pre-1.0 project, backward compatibility is *not guaranteed* – and indeed there have been breaking changes as features were added (for example, output formats and flags introduced in 0.4/0.5 could change scripting behavior). For production adoption, the team should:

- **Declare Stability at 1.0:** With the first major release, indicate that the CLI's interface (command names, flags, and output structure for machine-readable formats) will remain stable. Any breaking changes will be rare and tied to major version increments. This gives users confidence to write automation around RAPS. Adhering to Semantic Versioning strictly is critical – the project already signals intent to do so ⁵⁶.
- **Graceful Deprecation:** If any breaking change is needed post-1.0, use deprecation warnings. For example, if a flag is renamed, have the old flag still work but print a warning that it will be removed in a future release. This way users have time to adapt. Also, maintain the changelog diligently with a “**Deprecated**” or “**Breaking Changes**” section when applicable.
- **Testing Backward Compatibility:** As part of CI, it can be useful to have a few integration tests that use a known older version’s behavior to ensure it remains consistent. For instance, if JSON output for `raps bucket list` has a certain schema, write a test expecting that schema. This way, if a code change accidentally modifies the output structure, the test will catch it, flagging a potential breaking change. Essentially, treat the CLI output (at least in JSON/YAML modes) as an API contract.
- **Communication:** Clearly communicate in the README or docs that until 1.0, things may change, but after 1.0, compatibility will be a priority. Since many users will install via `cargo` or binaries without reading all commit notes, consider using the GitHub Releases to highlight any changes. The current changelog is very detailed, which is excellent; keep using that as the definitive reference for upgrades.

Given the thoroughness of the changelog and the semantic version approach, RAPS is on the right path. Locking down the interface and ensuring future changes don’t break existing scripts (or if they do, that users are well-warned) will be the final step in making it a **trustworthy production tool**.

Conclusion – Recommendations Summary: To reach production-grade quality, RAPS should focus on finishing incomplete features (ACC modules, plugin system) and continue polishing what’s already there. Key recommendations include improving pagination and performance for scale, enabling more secure defaults (use OS keychain by default, signed releases), expanding tests and CI to prevent regressions, and setting a clear path to a stable 1.0 release with strong backward compatibility. The good news is that as of v0.7.0, RAPS already has a solid foundation: it covers the core APS functionality with a user-friendly CLI, has extensive documentation and examples, and incorporates many best practices like standardized outputs and error codes for automation ³⁶ ⁶². By addressing the remaining gaps – both in features and in operational polish – the project can confidently call itself *production-ready*. The roadmap items in the pipeline (both literally and figuratively) should be completed and any known limitations transparently documented. With those improvements, RAPS will be well-positioned as a reliable, efficient, and secure APS command-line tool for both developers and enterprise CI workflows.

Sources:

- Project README and Documentation (features, installation, usage) 63 8
 - RAPS Roadmap and Implementation Status docs 5 17
 - Changelog entries for versions 0.4.0–0.7.0 64 38
 - Repository code and comments (configuration, plugin system, error handling) 58 3
 - Contribution and Release guidelines 52 54
 - GitHub issue tracker and planned enhancements 33 65
-

1 mod.rs

<https://github.com/dmytro-yemelianov/raps/blob/261375d5e29021e119b3b87b644c3d73a891c28f/src/api/mod.rs>

2 5 11 16 19 21 34 36 38 39 50 54 55 56 64 CHANGELOG.md

<https://github.com/dmytro-yemelianov/raps/blob/261375d5e29021e119b3b87b644c3d73a891c28f/CHANGELOG.md>

3 4 10 59 plugins.rs

<https://github.com/dmytro-yemelianov/raps/blob/261375d5e29021e119b3b87b644c3d73a891c28f/src/plugins.rs>

6 14 32 37 44 47 ROADMAP.md

<https://github.com/dmytro-yemelianov/raps/blob/261375d5e29021e119b3b87b644c3d73a891c28f/roadmap/ROADMAP.md>

7 17 18 20 22 23 24 25 26 27 28 29 40 41 48 62 65 IMPLEMENTATION_STATUS.md

<https://github.com/dmytro-yemelianov/raps/blob/261375d5e29021e119b3b87b644c3d73a891c28f/>

IMPLEMENTATION_STATUS.md

8 9 15 examples.md

<https://github.com/dmytro-yemelianov/raps/blob/261375d5e29021e119b3b87b644c3d73a891c28f/docs/examples.md>

12 13 features.md

<https://github.com/dmytro-yemelianov/raps/blob/261375d5e29021e119b3b87b644c3d73a891c28f/docs/features.md>

30 31 roadmap-v0.4-v0.6.json

<https://github.com/dmytro-yemelianov/raps/blob/261375d5e29021e119b3b87b644c3d73a891c28f/roadmap/roadmap-v0.4-v0.6.json>

33 Proxy support documentation ('HTTP_PROXY', 'HTTPS_PROXY', 'NO_PROXY')

<https://github.com/dmytro-yemelianov/raps/issues/52>

35 v0.7.0: Major feature release

<https://github.com/dmytro-yemelianov/raps/pull/60>

42 43 sign-binaries.ps1

<https://github.com/dmytro-yemelianov/raps/blob/261375d5e29021e119b3b87b644c3d73a891c28f/scripts/sign-binaries.ps1>

45 46 auth.md

<https://github.com/dmytro-yemelianov/raps/blob/261375d5e29021e119b3b87b644c3d73a891c28f/docs/commands/auth.md>

49 pipeline.md

<https://github.com/dmytro-yemelianov/raps/blob/261375d5e29021e119b3b87b644c3d73a891c28f/docs/commands/pipeline.md>

51 53 63 README.md

<https://github.com/dmytro-yemelianov/raps/blob/261375d5e29021e119b3b87b644c3d73a891c28f/README.md>

[52](#) CONTRIBUTING.md

<https://github.com/dmytro-yemelianov/raps/blob/261375d5e29021e119b3b87b644c3d73a891c28f/CONTRIBUTING.md>

[57](#) [58](#) [60](#) [61](#) configuration.md

<https://github.com/dmytro-yemelianov/raps/blob/261375d5e29021e119b3b87b644c3d73a891c28f/docs/configuration.md>